


EMU Chick Overview and Hands-On

31 July 2019

Hosted by Will Powell and Jason Riedy

Emu Technology (c) material will be denoted by the use of the Emu logo

CREATING THE NEXT MOORE'S LAW




Outline

- Introduction to the Emu Chick **(1:30-2:00 PM)**
 - Example 1: Hello World
 - Memory replication
 - Basic thread Spawning Strategies
- Programming for the Chick **(2:00-4:30 PM)**
 - Example 2: STREAM
 - Spawn strategies
 - Granularity
 - Locality awareness
 - Example 3: Sparse Matrix Vector Multiply (SpMV)
 - COO for fully interleaved operation
 - CSR decomposition for the Chick



Emu Chick Overview



Emu Innovation Overview

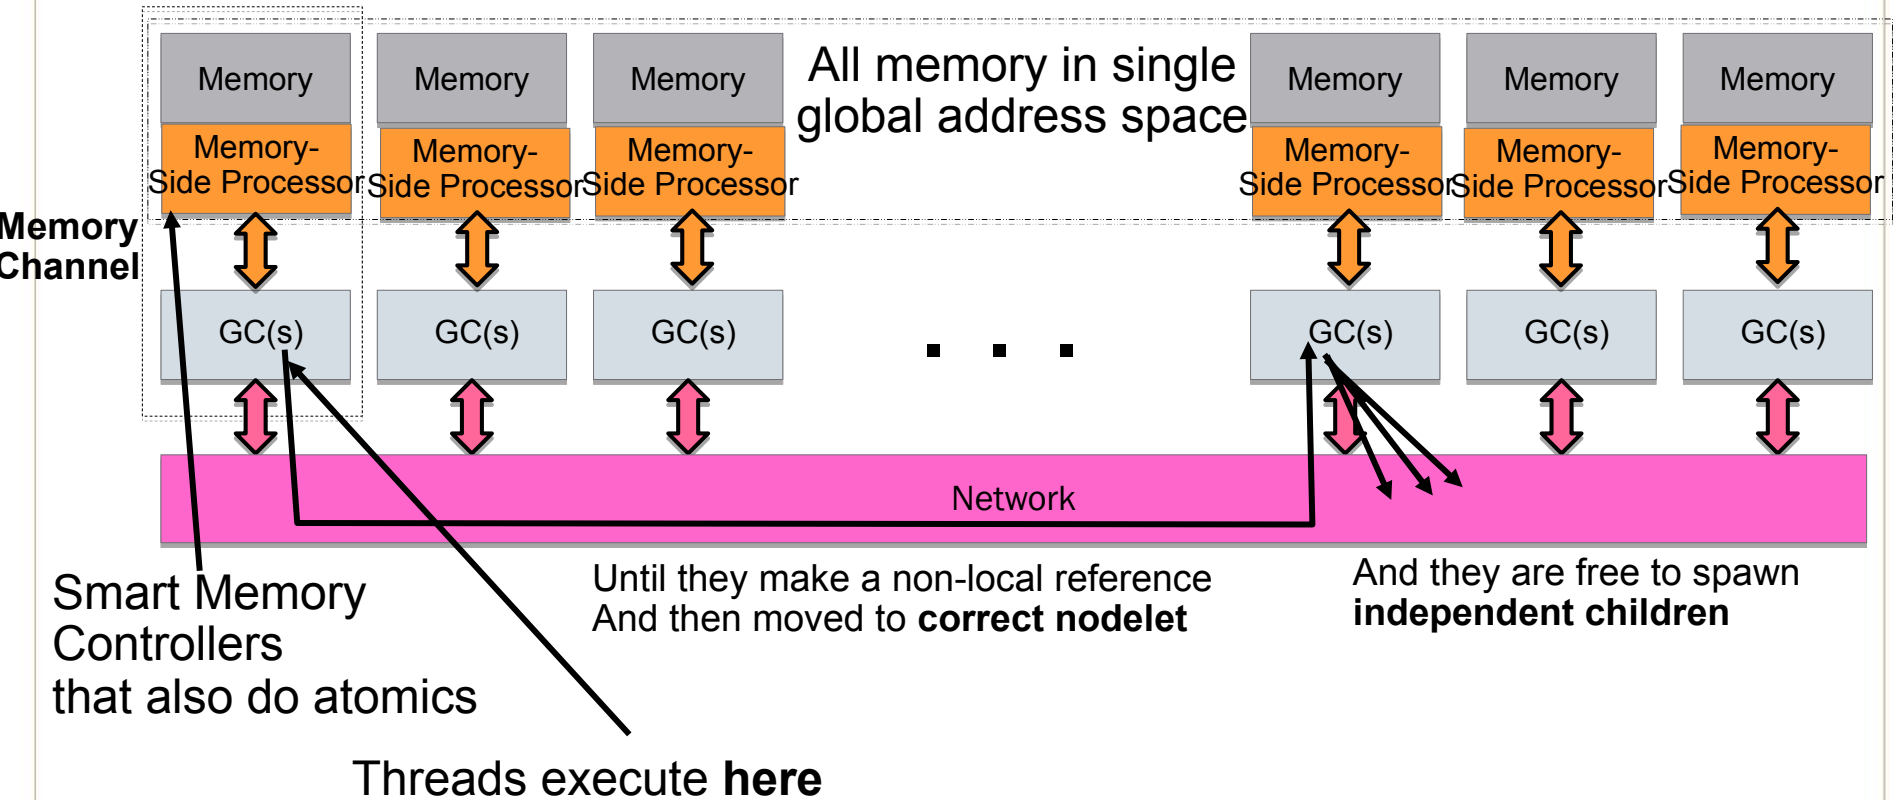
Designed from the ground up to deal with applications that exhibit little locality

- Massive Shared Memory for in-Memory Computing
 - No I/O bottlenecks
- **EMU** moves (“**Migrates**”) the program context to the locale of the data accessed
 - Lower energy – less data moved shorter distances
- Finely Grained Parallelism
 - Reduces concurrency limits
- Compute, memory size, memory bandwidth and software scale simultaneously

EmuTechnology

Emu Architecture Functional Diagram

Nodelet: New unit of parallelism




Node Architecture

- 8 Nodelets
- Migration Engine
- 6 RapidIO 2.3 4-lane network ports
- Stationary Cores (SCs)
DualCore 64-bit Power E5500
 - 2GB DRAM
 - 1 TB SSD
 - PCIe Gen 3
 - Runs Linux



Stationary Core
Runs OS, Launches
Jobs




Migrating Threads
are major traffic
on Network

Nodelet Architecture




- 8 GB DDR4 Narrow Channel Memory
 - Supports 64-bit accesses
- Memory-side Processor (MSP)
 - Handles atomics and remote writes at the memory
- Gossamer Cores (GCs) each with FMA FPU
- Nodelet Queue Manager
 - Run Queue
 - Incoming threads from migrations, spawns, or SC
 - Loaded into vacant execution slots by hardware
 - Migration Queue
 - Threads that need to migrate to non-local data
 - Service Queue
 - Threads that need system services from the SC



(a) Nodelet

**Multi-Threaded
Cores**




Hardware Thread Management

- Thread scheduling in GCs automatically performed by hardware
- SPAWN instruction
 - Creates new thread and places it in Run Queue
- RELEASE instruction
 - Places thread in Service Queue for processing by SC
- Non-local memory reference causes a migration
 - Thread context packaged by hardware and placed in Migration Queue
 - Migration Engine sends packet to new location and places in Run Queue


Emu System Hierarchy

Atomics run
in Memory-Side
Processor (MSP)



(a) Nodelet


**Multi-Threaded
Cores**



(b) Node

**Migrating Threads
are major traffic
on Network**

**Stationary Core
Runs OS, Launches
Jobs**




Up Links to
(c) Inter-Chassis
Switch



Emu Chick Topology

- System consists of 8 nodes connected in a cube via RapidIO links
 - Each node connects to 6 other nodes
 - Cube edges and face diagonals are connected, but not interior diagonals
- All routes are 2 hops or 1
 - 3D diagonals route through intermediate node
 - All others are 1 hop




The Current Chick Hardware





- 8 nodes (64 nodelets)
- 512 GB Shared Memory
- 8 TB SSD
- 8192 Concurrent Threads
- Copy room environment
- Shipping now





Emu Chick Architecture






Emu's Migratory Thread Model

Massive, fine-grained multithreading where computation migrates to the data so that accesses are always local


Key Issues:

- Thread control: spawning and synchronization
- Data distribution and affinity of execution
 - Load balance
 - Hotspots
 - Migration patterns




Fine-grained Memory Accesses

- Narrow-channel DRAM (NCDRAM)
 - 8-bit bus allows access at 8-byte granularity without waste
 - Many narrow channels instead of few wide channels
- Remote Writes
 - Write to remote nodelet without migrating
 - Proceed directly to the memory front-end, bypassing the GC
- Remote Atomics
 - Performed in Memory Front-End (MFE), near memory



Emu Programming – Key Features

- **Cilk:** Extensions to C to support thread management
 - cilk_spawn
 - cilk_sync
 - cilk_for
- **Emu Cilk:** Extensions to Cilk to support migrating threads
 - cilk_migrate_hint
 - cilk_spawn_at



Emu Programming – Key Features


- **Memory allocation library:**
Specialized malloc/free for data distributed across nodelets
- **Intrinsics:** Allow access to architecture specific operations such as atomic updates



Emu Cilk

Emu hardware dynamically creates and schedules threads


- Normally requires no software intervention
- When a thread completes, it returns values to its parent and dies
- When a thread blocks, it may voluntarily place itself at the back of the run queue (instead of “busy waiting”)
- Number of threads limited only by available memory
- Extremely lightweight – Cilk threads can be very small and still be efficient



Cilk Functions

```
long f = cilk_spawn fib(a, b);
```


- Specifies function may run in parallel with caller
 - Child thread spawned to execute function and parent continues in parallel w/child
 - Otherwise parent executes a standard function call
- Spawn location determines location of
 - Synchronization structure
 - Stack frame (if needed)
- Spawn destination
 - Special functions denote spawn location
 - If no direction is given, then spawn is local



Cilk Functions

`cilk_sync`;


- Current function cannot continue past the `cilk_sync` until all children have completed
- Last thread to reach the `cilk_sync` continues execution – **no waiting**
- Implicit sync at termination of a function



Cilk Functions

```
#pragma cilk grainsize = 4  
cilk_for(long i=0; i<SIZE; i++)  
    {...}
```

- Divides loop among parallel threads, each containing one or more contiguous loop iterations
- Max number of iterations in each chunk is *grainsize*
- Best for situations where
 - Threads are spawned locally
 - Work per element is fairly uniform



Emu CilkPlus Functions

cilk_migrate_hint(p);

- Specifies nodelet for next cilk_spawn operation
 - Argument p is a pointer into destination nodelet's memory

cilk_spawn_at(p) fib(a,b);

- Combines cilk_migrate_hint and cilk_spawn into a macro for single-command spawn
 - Implemented as C macro; may require braces for correct operation



Cilk Fibonacci Example


```
1  #include "memweb.h"
2  #include <cilk/cilk.h>
3  #define N 10
4  long fib(long n) {
5      if (n < 2)
6          return n;
7      long a = cilk_spawn fib(n-1);
8      long b = cilk_spawn fib(n-2);
9      cilk_sync;
10     return a + b;
11 }
12 int main() {
13     long result = fib(N);
14     printf("fib(%d) = %ld\n", N, result);
15 }
```

Spawn a thread
for each of the
fib() calls

Wait for threads to
complete to ensure
a and b are valid



HELLO WORLD AND HANDS-ON



Ex 1: Emu Hello World


This example demonstrates the following:

- Memory replication
- Basic thread Spawning Strategies

Chick Hands-On Information



- System: **notebook.crnch.gatech.edu**
- Username: pearc\$(seq 0 49)
- Login: <https://notebook.crnch.gatech.edu>
Jupyterhub → ephemeral JupyterLab
- Emu directory: /tools/emu
- Already in \$PATH:
 - Compiler: emu-cc [-h]
 - Simulator: emusim.x [-h]




Emusim.x can be used to explore differences in..

- Execution Time
- Migrations
- Memory Map
- Remotes Map
- Run Queue
- Migration Queue
- Remote Queue

Configuration and Summary Statistics

- Generated automatically in `<program>.cdc`
- Check number of threads spawned, their distribution, and number of migrations
- View memory map and remotes map to identify hotspots, poor distribution, and migration patterns




Examine Simulation Output

- `hello_world.cdc`
 - Execution Time
 - Total Threads
 - Maximum Concurrent Threads
 - Memory Map
 - Remotes Map




Verbose Simulation Statistics

- Generated automatically in <program>.vsf
- Identify hotspots/bottlenecks by examining max queue depths
 - Run Queue
 - Migration Queue
 - Remote Queue
- Identify utilization at nodelets
 - IPC (Instructions per cycle: Max 4.0 for 4 cores)
 - Memory Bandwidth (Max 1.0)
 - System IC Bandwidth (Max 1.0)




Examine Simulation Output


- `hello-world.vsf`
 - Much of the same information presented with more detail and different organization
 - Provides queue depths for run queue, migration queue, and remote queue
- Run visualization tool
`emuvistool hello-world.tqd`



Ex 1: Exploring Simulation Output



EMU MEMORY ALLOCATION



Memory Allocation

- Replicated, Stack, and Heap sections on each nodelet
- Replicated – global replicated data
- Stack – local memory allocation
 - Thread frames
 - `malloc()/free()`
 - `new()/delete()`
- Heap – distributed memory allocation
 - Specialized `mw_*malloc*()` functions



Global Replicated Data Structures

replicated long c = 3927883;

- Instructs compiler to place an instance on each nodelet
- Uses a “View 0” address that always gives local instance
- Must be a **global** variable
- Example Uses:
 - Constants
 - Copy on each nodelet
 - All initialized to the **same** unchanging value
 - EX: PI, pointer to shared data structure
 - Local data
 - Copy on each nodelet
 - May have **different** values
 - Use only when it does **not** matter which instance you access!
 - EX: random number table, pointer to local work queue

Dynamic Replicated Pointers

`long * mw_mallocrepl(size_t blocksize)`

- Allocates a block on each nodelet, returns replicated pointer
- Similar to using the replicated keyword
- Used when the size of the data structure is not known at compile time

Replicated Data Structures

- Replicating key shared data structures can improve performance
 - Pointers to shared distributed data e.g. array
 - Copy at each nodelet avoids migrations to get address
 - Compiler generates the address rather than having to pass the address to each function call and carry it during migrations
 - Can reduce spills at function calls



Initializing Replicated Data Structures

```
void mw_replicated_init(long *repl_addr, long value)
```

- Initializes each instance of replicated data structure to value

```
void mw_replicated_init_multiple (long *repl_addr,  
                                  long (*init_func)(long) )
```

- Initializes each instance of replicated data structure using the **result** of the user-defined function `init_func(n)` where `n` is the nodelet number

```
void mw_replicated_init_generic(long *repl_addr,  
                                void (*init_func)(void *, long) )
```

- Initializes each instance of replicated data structure using the user-defined function `init_func(&obj, n)`, where `obj` is the address of the replicated data structure and `n` is the nodelet




Accessing Replicated Data Structures

```
void * mw_get_localto(void *r_ptr, void *dest_ptr)
```

- Returns a pointer to the instance of a replicated data structure co-located with the destination pointer


```
void * mw_get_nth(void *r_ptr, unsigned n)
```

- Returns a pointer to the nth instance of a replicated data structure



Local Memory Allocation

- Allocate from the stack on the current nodelet using conventional C/C++ functions
 - malloc and free
 - new and delete



Distributed Memory Allocation

`void * mw_localmalloc(size_t eltsize, void *ptr)`

- Block of memory located in same locale as another data structure


`void * mw_malloc1dlong(unsigned numelements)`

- Array of longs striped across nodelets round robin

`void ** mw_malloc2d(unsigned nelements,
size_t eltsize)`

- Array of pointers striped across nodelets round robin
- Each points to a block of memory in the same locale

Distributed Free



void mw_free(void *allocatedpointer)

- Free data allocated by mw_malloc2d

void mw_localfree(void *allocatedpointer)

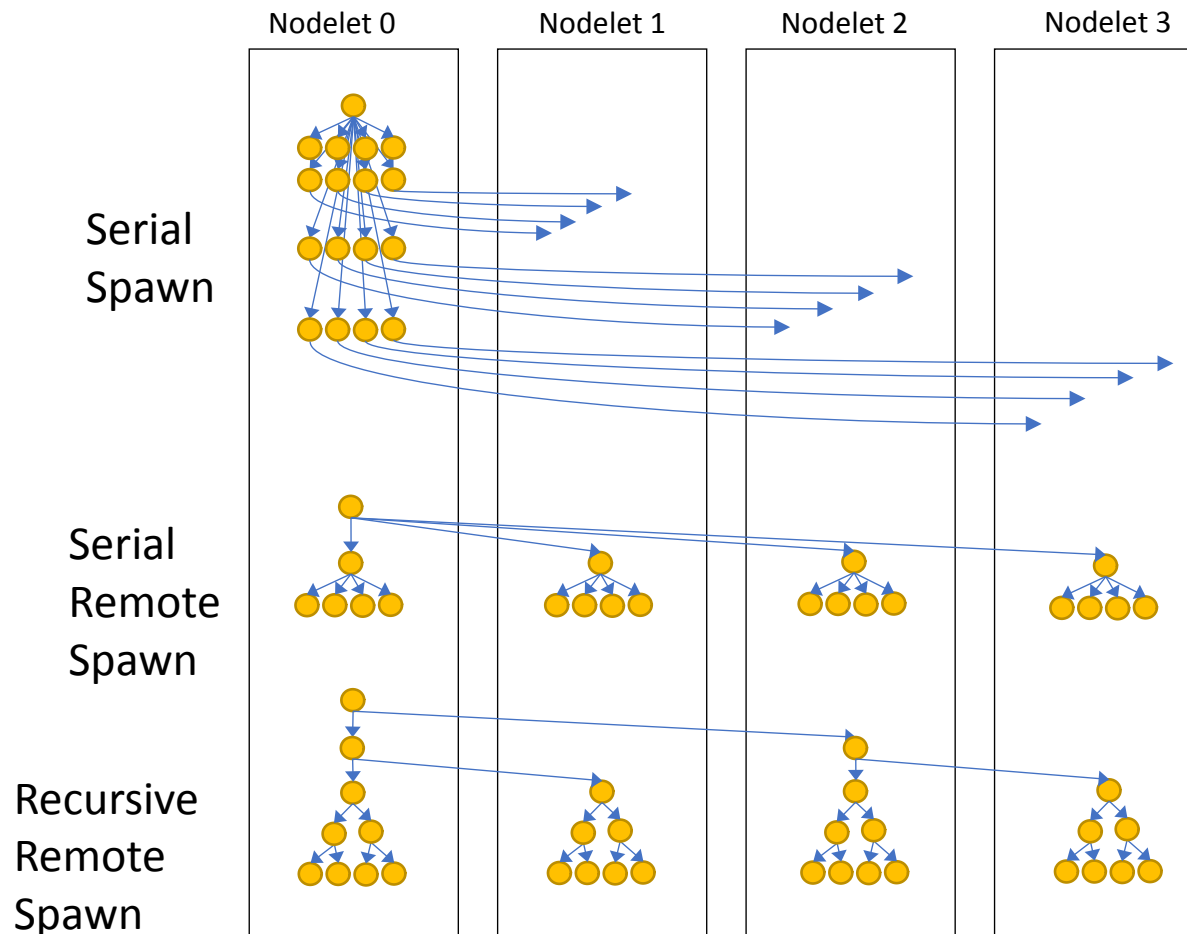
- Free data allocated by mw_localmalloc


Accessing Distributed Data

```
long * mw_arrayindex(void *array2d,  
unsigned long i, unsigned long numblocks,  
size_t blocksize)
```

- Inputs:
 - Array allocated with mw_malloc2d
 - Index for first dimension
 - Number of blocks used in malloc2d
 - Blocksize used in malloc2d
- Returns address of array2d[i][0]

Added Spawn Strategies






STREAM Implementation


```
// Serial
for (long i = 0; i < n; ++i) {
    c[i] = a[i] + b[i];
}

// Parallel
cilk_for (long i = 0; i < n; ++i) {
    c[i] = a[i] + b[i];
}
```




Ex: 2 Emu STREAM

- This example demonstrates the following:
- More complex spawning strategies
 - Grain size settings
 - Locality awareness (i.e., limiting some migrations)






HARDWARE EXECUTION OVERVIEW



Emu Chick Configuration


- Single-node Execution
 - Program runs on a single node
 - Can access all 8 nodelets but no other nodes
 - Users can work independently on different nodes
- Multi-node Execution
 - Program runs on full system
 - Can access all 8 nodes (64 nodelets)
 - Single user




Emu Chick Hardware Execution

- Compile programs on notebook.crnch.gatech.edu then scp to Chick node
- Single-node Execution
 - Launched on node using `emu_handler_and_loader`
- Multi-node Execution
 - Launched on node 0 using `emu_multinode_exec`

Program Execution Utilities




- **Load** program and data to all nodelets
- **Launch** initial thread into the system
- **Monitor** the system exception queue and handle system services until a thread quits or an exception occurs
- **Terminate** by issuing a checkpoint to clear the system and dump any remaining threads
- **Print** information to log files for each thread that quits, exits, generates an exception, or is checkpointed
- **Return** the program's return value



System Services

- Thread suspends itself and writes thread state registers (TSR) to the system exception queue (SEQ)
- Handler polls system SEQ for threads that need services
- Handler reads the thread from the SEQ and performs the requested service
- Handler then relaunched the thread to nodelet 0


Running Code on the Chick



- System: **karrawingi-login.crnch**
 - SSH keys should be set up for you to use
 - `~/.ssh/config` includes proxy setup
- Username: `pearc$(seq 0 49)`
 - *ssh karrawingi-login*
- Emu Chick nodes
 - *ssh n0 - n7*

Ex 3: Emu SpMV

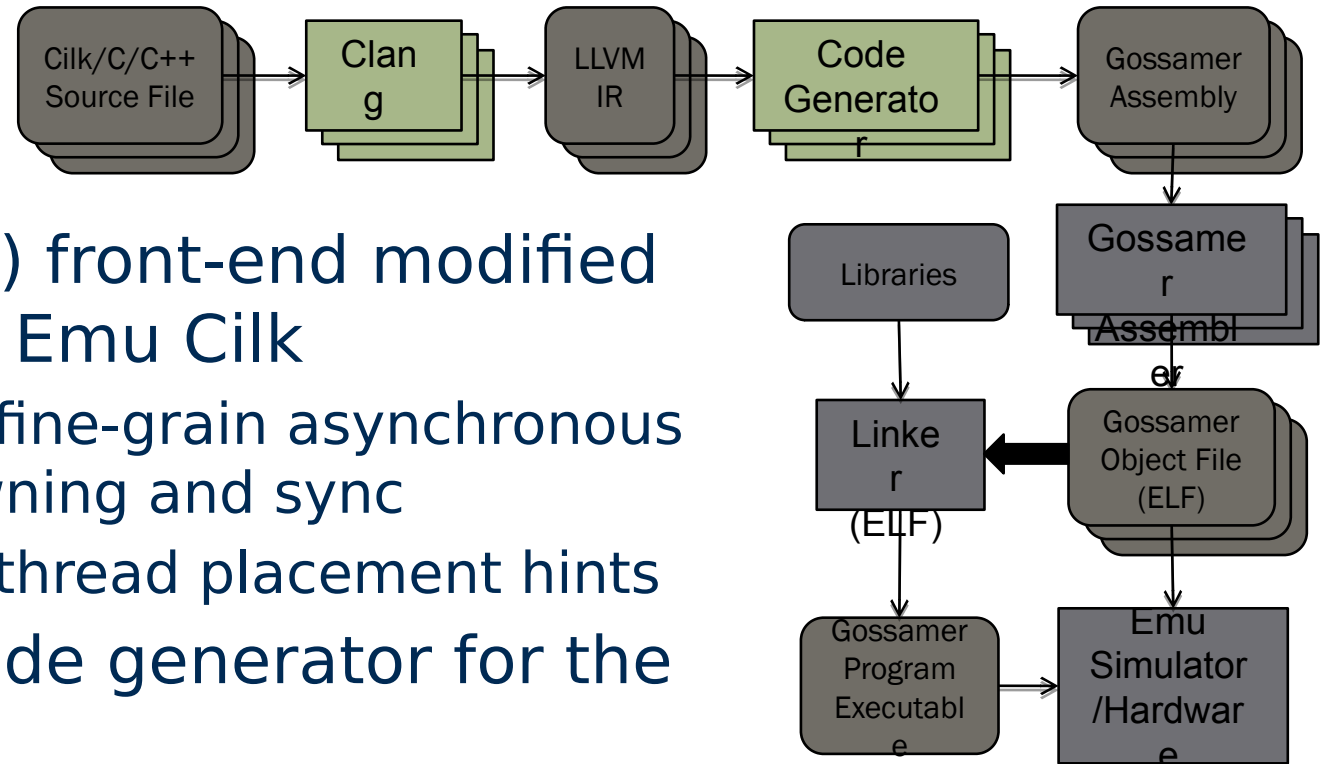
- This example demonstrates the following:
- COO vs CSR layouts
 - More advanced data distributions






EMU TOOLCHAINS AND DEVELOPMENT ENVIRONMENT

Emu Cilk Toolchain




- Cilk (clang) front-end modified to support Emu Cilk
 - Supports fine-grain asynchronous task spawning and sync
 - Supports thread placement hints
- Custom code generator for the Emu GCs
- Custom calling convention and run-time support
- Custom assembler and linker

Support for C, C++, and CilkPlus provides **familiar development environment**




CilkPlus Beta Release

- Latest Clang front-end to support
 - C/C++ 2011
 - CilkPlus
- Key CilkPlus features
 - **Reducers:** list, min/max, addition, bitwise AND/OR/XOR, multiplication, ostream, string, vector
 - **Pedigrees:** unique naming convention for threads
- No support for
 - CilkPlus vector operations
- Currently delivered in parallel with standard release




Emu C Utilities

- Set of common patterns for thread-parallel code implemented efficiently as library calls
- Working with local arrays
 - Alternative to `cilk_for`, no compiler support
- Working with distributed striped arrays
 - 2-level spawn tree, split array for worker functions
- Working with distributed chunked arrays
 - Calculates indices, applies functions to blocked arrays
- Timing hooks
 - Timer subsystem for performance analysis




User Libraries

- GNU Multiple Precision Arithmetic (GMP) Library
 - Library for arbitrary precision arithmetic
 - Currently support integer GMP for Emu
 - Included in current release
- Under development
 - GraphBLAS (UMBC / SEI)
 - OpenMP (Stony Brook University)
- Other efforts
 - STINGER Graph Library (Georgia Tech)
 - Kokkos C++ Ecosystem (Georgia Tech / Sandia)




What have we not covered today?

- Atomics and intrinsics
- Multi-node execution
- Thread management




Backup



Intrinsics

- Set of compiler recognized functions to access architecture specific operations
 - Atomic Arithmetic Operations
 - Remote Arithmetic Operations
 - Other Architecture Specific Operations
 - Thread Management Functions
 - System Queries



Atomic Arithmetic Operations

- Atomic arithmetic operations at the memory

```
long ATOMIC_<INST> (volatile void *A, long D);  
<INST>: ADD/SUB/AND/OR/XOR/MAX/MIN
```


- Four distinct “flavors” of each atomic that specify the value returned and resulting mem

Mnemonic	Return Value	New Memory Value
<op>	Result	Unchanged
<op>S	Result	Orig. D Value
<op>M	Result	Result
<op>MS	Orig. Mem. Value	Result



ATOMIC_ADD Examples


- **ATOMIC_ADD**(A, D)
 - Reads value at A
 - Adds D
 - Returns result
 - Mem[A] is unchanged
- **ATOMIC_ADDS**(A, D)
 - Reads value at A
 - Adds D
 - Returns results
 - Mem[A] = D
- **ATOMIC_ADDM**(A, D)
 - Reads value at A
 - Adds D
 - Returns result
 - Mem[A] = result
- **ATOMIC_ADDMS**(A, D)
 - Reads value at A
 - Adds D
 - Returns original Mem[A]
 - Mem[A] = result



Remote Arithmetic Operations


```
void REMOTE_<INST>(volatile void *A, long D);  
    <INST>:  ADD/AND/OR/XOR/MAX/MIN
```

- Significantly reduces migrations by performing atomic updates to memory *without migrating the thread*
- Sends only the data and operation to be performed
 - Consumes less than half the bandwidth of a typical thread migration
- Does not return a value
- Returns an ACK, may be turned off
- Remotes issued by a thread to the same location guaranteed to complete in order



WRD – Remote Write


- Automatically generated by the compiler in place of the store instruction
- Writes the value to remote memory and sends an acknowledgement (ACK)
- If the access is local, treated as a store



FENCE on Remote Update

```
void FENCE();
```

- Used to wait for all remote updates to be acknowledged
- Prevents thread from continuing until all outstanding acknowledgements have been received
- Implicit FENCE before migration, thread suspend, etc.




Swap Operations

long **ATOMIC_SWAP**(volatile void *A, long D)

- Replace the contents of memory at A with D
- Return the original contents of memory at A

long **ATOMIC_CAS**(volatile void *A,
long newVal, long cmpVal)

- Compare the contents of memory at A to cmpVal.
- If they match, swap newVal into A
- Always return original contents of memory at A



Other Operations


long **POPCNT**(long sum, long * ptr)

- Counts number of 1s in the word referenced by ptr
- Adds this value to sum and returns the result

unsigned long

PRIORITY(unsigned long value)

- Computes the 6 bit priority encode on value
- i.e. bit position of highest numbered non-



Thread Management

void **MIGRATE**(volatile void *A)


- Force a migration to an address

void **RESIZE**()

- Resize the thread to carry only the live registers
- Used by compiler or programmer before a possible migration to reduce thread size


void **RESCHEDULE**()

- Place the thread at the end of the Run Queue to allow a new thread to be scheduled in the core
- Can be used to minimize the impact of busy wait



Thread Management (cont)

- void `ENABLE_ACKS()` – require acks on remotes
- void `DISABLE_ACKS()` – turn off acks
- void `DISABLE_MIGRATIONS()` – cause exception on migration (for debugging)
- void `ENABLE_MIGRATIONS()` – resume migrations without exceptions
- void `DISABLE_INTERRUPTS()` – prevent thread from timeslice interrupt
- void `ENABLE_INTERRUPTS()` – allow timeslice interrupts from thread
- void `ENTER_CRITICAL_SECTION()` – combine `DISABLE_MIGRATIONS` and `DISABLE_INTERRUPTS`
- void `EXIT_CRITICAL_SECTION()` – combine `ENABLE_MIGRATIONS` and `ENABLE_INTERRUPTS`



System Query

- **CLOCK()** - Returns system time on local nodelet (clock ticks)
- **THREAD_ID()** - Returns the current thread ID
- **NODE_ID()** - Returns the current Nodelet ID
- **NODELETS()** - Returns the number of nodelets in the current system
- **BYTES_PER_NODELET()** - Returns the number of bytes per nodelet
- **MAXDEPTH()** - maximum spawn depth*



Initializing Replicated Data Structures

```
void mw_replicated_init(long *repl_addr, long value)
```

- Initializes each instance of replicated data structure to the same value

```
replicated long N;  
int main()  
{  
    ...  
    long n_elements = compute_n_elements();  
    mw_replicated_init(&N, n_elements);  
    ...  
}
```



Initializing Replicated Data Structures

```
void mw_replicated_init_multiple
```

```
    (long *repl_addr, long (*init_func)(long))
```

- Initializes each instance of the replicated data structure using the **result** of the user-defined function `init_func(n)` where `n` is the nodelet number

```
    replicated long B;
```

```
    long init_func(long nid) {
```

```
        return nid * 5;
```

```
    }
```

```
    int main()
```

```
    {    ...
```

```
        mw_replicated_init_multiple(&B, init_func);
```

```
        ...
```

```
    }
```





Initializing Replicated Data Structures

```
void mw_replicated_init_generic(long *repl_addr,  
                                void (*init_func)(void *, long) )
```

- Initializes each instance of replicated data structure using the user-defined function `init_func(&obj, n)` where `obj` is the address of the replicated data structure and `n` is the nodelet number


```
replicated struct info { long x; long y; } info;  
void init_info(void *obj; long nid) {  
    struct info *i = (struct info*) obj;  
    i->x = node;  
    i->y = 5*node + 4;  
}  
int main()  
{  
    ...  
    mw_replicated_init_multiple(&info, init_info);  
    ...  
}
```




Allocate Replicated Data Structure

Copy of stats on each nodelet at the same offset


```
struct stats {  
    unsigned long count;  
    unsigned long max;  
    unsigned long min;  
};  
replicated struct stats s;
```





Initialize Replicated Data Structure


```
// Initialize using mw_get_nth  
for (long i=0; i<NODELETS(); i++) {  
    struct stats * si = mw_get_nth(&s, i);  
    si->count = 0;  
    si->max = 0;  
    si->min = LONG_MAX;  
}
```



Update Replicated Data Structure


```
// Update stats
cilk_for (long i=0; i<N; i++) {
    unsigned long score = score(A[i]);
    if (score != 0) {
        ATOMIC_ADDM(&(s.count), 1);
        ATOMIC_MAXM(&(s.max), score);
        ATOMIC_MINM(&(s.min), score);
    }
}
```


Remember
to use
atomics for
shared data



Reduce Replicated Data Structure

```
// Reduce using mw_get_nth
unsigned long count = 0;
unsigned long max = 0;
unsigned long min = LONG_MAX;
for (long i=0; i<N; i++) {
    struct stats * si = mw_get_nth(&s, i);
    count += si->count;
    if (si->max > max) max = si->max;
    if (si->min < min) min = si->min;
}
```







Allocate Co-located Data

```
void * mw_localmalloc(size_t eltsize, void *ptr)
```

- Block of memory of size `eltsize` co-located with address `ptr`
- Accessed using traditional C semantics
- Leaves you at the new location
- Example:

```
long * A = (long *)  
    mw_localmalloc(sizeof(long), &B[2]);
```






Allocate 1D Array of Longs

```
long * mw_malloc1dlong(unsigned nelements)
```


- Array of nelements longs striped across nodelets round robin
- Accessed using 1D array notation (e.g. A[i])
- ONLY works for 64-bit types (e.g. long)




Example: 1D Array

Array of N elements, striped across nodelets round-robin

```
#define N 16  
long * A = (long *) mw_malloc1dlong(N);  
for (long i=0; i<N; i++)  
    A[i] = i;
```






Allocate 2D Array of Elements


```
void ** mw_malloc2d(unsigned nelements,  
                    size_t eltsize)
```

- Array of nelements pointers striped across nodelets round robin
- Each points to co-located memory block of size eltsize
- May be an array of pointers to any type



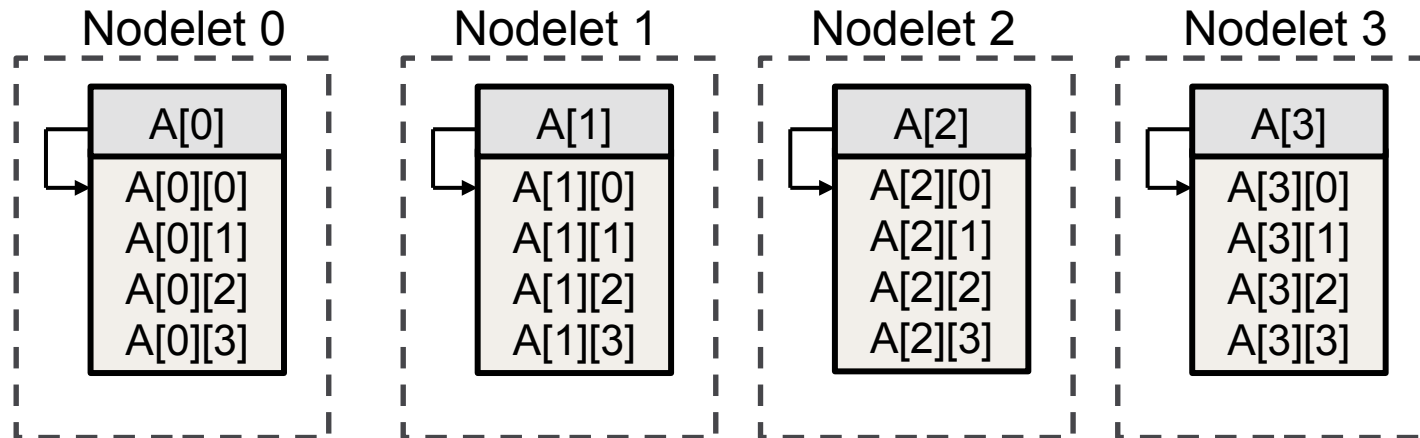
Malloc2d Ex: Array of structs

```
#define N 8
struct element { long a; long b; };
struct element ** E = (struct element **)
    mw_malloc2d(N, sizeof(struct element));
for (long i=0; i<N; i++) {
    E[i]->a = i;
    E[i]->b = 0;
}
```



Malloc2d Ex: Blocked Distribution

```
#define N 16 // # elements in array (power of 2)
long epn = N/NODELETS(); // elements per nodelet
long ** A = (long **)
    mw_malloc2d(NODELETS(), epn * sizeof(long));
for (long i=0; i<N; i++)
    A[i/epn][i%epn] = i;
```



Malloc2d Ex: Wrapped Block Dist.

```
#define N 16
long nBlocks = 8;
long blockSize = N/nBlocks;
long ** A = (long **)
    mw_malloc2d(nBlocks, blockSize * sizeof(long));
for (long i=0; i<N; i++)
    A[i/blockSize][i%blockSize] = i;
```

